

ECE 565 HW5 Report

Sequential:

Data structure:

In this problem we used mainly arrays to store different variables.

```
int timeSteps = stoi(argv[1]);
float absRate = stof(argv[2]);
int N = stoi(argv[3]);
string elevation_file = argv[4];
float runtime;
vector<vector<int>> elevation(N, vector<int>());
vector<vector<float>> absorb(N, vector<float>(N, 0));
struct timespec start_time, end_time;
int wholeSteps = 0; // Store the whole timesteps
```

```
int N = elevation.size();
vector<vector<float>> rain(
    N, vector<float>(N, 0)); // Store the current rain on the ground
vector<vector<float>> trickle(
    N, vector<float>(N, 0)); // Store the trickle of each step
vector<vector<float>> nextTrickle(
    N,
    vector<float>(N, 0)); // Store the trickle result to be added next round
vector<vector<float>> tempTrickle(
    N,
    vector<float>(N, 0)); // Store the trickle result to be added next round
vector<vector<float>> resetTrickle(
    N, vector<float>(N, 0)); // Used for resetting the tempTrickle
vector<vector<vector<vector<int>>>> neighborsToTrickle(
    N, vector<vector<vector<int>>>(N, vector<vector<int>>()));
```

Algorithm:

1. In main function:

First we initialize all the variables(elevation, absorb, wholeSteps...):

```
// Initialization of variables
int timeSteps = stoi(argv[1]);
float absRate = stof(argv[2]);
int N = stoi(argv[3]);
string elevation_file = argv[4];
float runtime;
vector<vector<int>> elevation(N, vector<int>());
vector<vector<float>> absorb(N, vector<float>(N, 0));
```

```

struct timespec start_time, end_time;
int wholeSteps = 0; // Store the whole timesteps

// open the file to read.
fstream in("./" + elevation_file);
string line;
int ele; // Store elevation for each input unit
if (in) // file exists
{
    int row = 0;
    while (getline(in, line)) // line does not contain newline of each line
    {
        stringstream ss(line);
        while (ss >> ele) {
            elevation[row].push_back(ele);
        }
        ++row;
    }
} else // No such file
{
    cout << "No such elevation_file" << endl;
    return EXIT_FAILURE;
}

```

Then we start calculating time and the steps needed to fully absorb all the rain:

```

wholeSteps =
    calcRain(elevation, absorb, timeSteps, absRate, start_time, end_time);

float elapsed_ns = calc_time(start_time, end_time);
cout << "Rainfall simulation took " << wholeSteps
    << " time steps to complete." << endl;
cout << "Runtime = " << elapsed_ns / 1000000000 << " seconds" << endl;
cout << endl;
cout << "The following grid shows the number of raindrops absorbed at each "
    "point:"
    << endl;

```

2. calcRain function:

In each iteration, we call “**rainAbsorbTrickle**” to receive new raindrops, absorb and trickle the water until the ground is fully dry, then return the wholesteps.

```

float isDrain = 1;

```

```

clock_gettime(CLOCK_MONOTONIC, &start_time);

while (isDrain != 0) {
    isDrain = 0;

    // 1) Receive a new raindrop (if it is still raining) for each point.
    // 2) If there are raindrops on a point, absorb water into the point
    // 3a) Calculate the number of raindrops that will next trickle to the
    // lowest neighbor(s)
    rainAbsorbTrickle(rain, absorb, trickle, timeSteps, absRate, isDrain,
                     nextTrickle, elevation, neighborsToTrickle, tempTrickle);
    nextTrickle = tempTrickle;
    tempTrickle = resetTrickle;
    --timeSteps;
    ++wholeSteps;
}

clock_gettime(CLOCK_MONOTONIC, &end_time);

return wholeSteps;

```

Here some some variables we pass into “rainAbsorbTrickle”:

```

int N = elevation.size();
vector<vector<float>> rain(
    N, vector<float>(N, 0)); // Store the current rain on the ground
vector<vector<float>> trickle(
    N, vector<float>(N, 0)); // Store the trickle of each step
vector<vector<float>> nextTrickle(
    N,
    vector<float>(N, 0)); // Store the trickle result to be added next round
vector<vector<float>> tempTrickle(
    N,
    vector<float>(N, 0)); // Store the trickle result to be added next round
vector<vector<float>> resetTrickle(
    N, vector<float>(N, 0)); // Used for resetting the tempTrickle
vector<vector<vector<vector<int>>>> neighborsToTrickle(
    N, vector<vector<vector<int>>>(N, vector<vector<int>>()));
for (int i = 0; i < N; ++i) {
    for (int j = 0; j < N; ++j) {
        vector<vector<int>> neighToTrickle = countNeighbor(i, j, elevation);
        neighborsToTrickle[i][j] = neighToTrickle;
    }
}

```

```
}
```

“countNeighbor” is used to sort and determine which neighbor(s) to trickle:

```
vector<vector<int>> countNeighbor(int i, int j,
                                const vector<vector<int>> &elevation) {
    int N = elevation.size();
    vector<vector<int>> neighToTrickle; // <elevation, number>, number: Left: 0,
                                        // up: 1, right: 2, down: 3
    vector<vector<int>> allNeigh;       // All neighbors that have less elevation
    if (j - 1 >= 0 && elevation[i][j - 1] < elevation[i][j]) { // left
        allNeigh.push_back({elevation[i][j - 1], 0});
    }
    if (i - 1 >= 0 && elevation[i - 1][j] < elevation[i][j]) { // up
        allNeigh.push_back({elevation[i - 1][j], 1});
    }
    if (j + 1 < N && elevation[i][j + 1] < elevation[i][j]) { // right
        allNeigh.push_back({elevation[i][j + 1], 2});
    }
    if (i + 1 < N && elevation[i + 1][j] < elevation[i][j]) { // down
        allNeigh.push_back({elevation[i + 1][j], 3});
    }
}
```

3. rainAbsorbTrickle function:

- 3.1 first we add the trickle calculated from last iteration, then add the new raindrop to absorb:

```
// Add trickle from the previous step
rain[i][j] += nextTrickle[i][j];
// Reset the nextTrickle array
nextTrickle[i][j] = 0;
if (timeSteps > 0) {
    // 1) Receive a new raindrop (if it is still raining) for each point.
    ++rain[i][j];
}
if (rain[i][j] == 0) {
    continue;
}
// 2) If there are raindrops on a point, absorb water into the point
if (rain[i][j] >= absRate) {
    rain[i][j] -= absRate;
    absorb[i][j] += absRate;
} else if (rain[i][j] > 0) {
    absorb[i][j] += rain[i][j];
    rain[i][j] = 0;
}
```

```

        continue;
    }

```

- 3.2 Then we calculate the amount of rain that will trickle to the near lowest neighbor(s).

```

// 3a) Calculate the number of raindrops that will next trickle to the
// lowest neighbor(s)
trickle[i][j] = 0; // Reset the trickle array
if (rain[i][j] >= 1) {
    trickle[i][j] = 1;
} else if (rain[i][j] > 0) {
    trickle[i][j] = rain[i][j];
}
isDrain += trickle[i][j];

```

(In here, we use **isDrain** to help check if the all the ground is dry)

- 3.3 For each point, we use the calculated trickle amount to update the actual raindrop that will trickle to the neighbor(s), if applicable.

```

calcTrickle(i, j, rain, elevation, trickle, neighborsToTrickle,
            tempTrickle);

```

4. calcTrickle function

In this function, we trickle the rain on each point based on the evaluation of their neighbors, and store the results in array **tempTrickle**, this will be used to update the **nextTrickle** array, which is used at the beginning of **rainAbsorbTrickle** to absorb the trickled water.

```

void calcTrickle(int i, int j, vector<vector<float>> &rain,
                const vector<vector<int>> &elevation,
                vector<vector<float>> &trickle,
                vector<vector<vector<vector<int>>>> &neighborsToTrickle,
                vector<vector<float>> &tempTrickle) {
    if (trickle[i][j] == 0 || neighborsToTrickle[i][j].size() == 0) {
        return;
    }
    vector<vector<int>> direct = {
        {0, -1}, {-1, 0}, {0, 1}, {1, 0}}; // Left, up, right, down

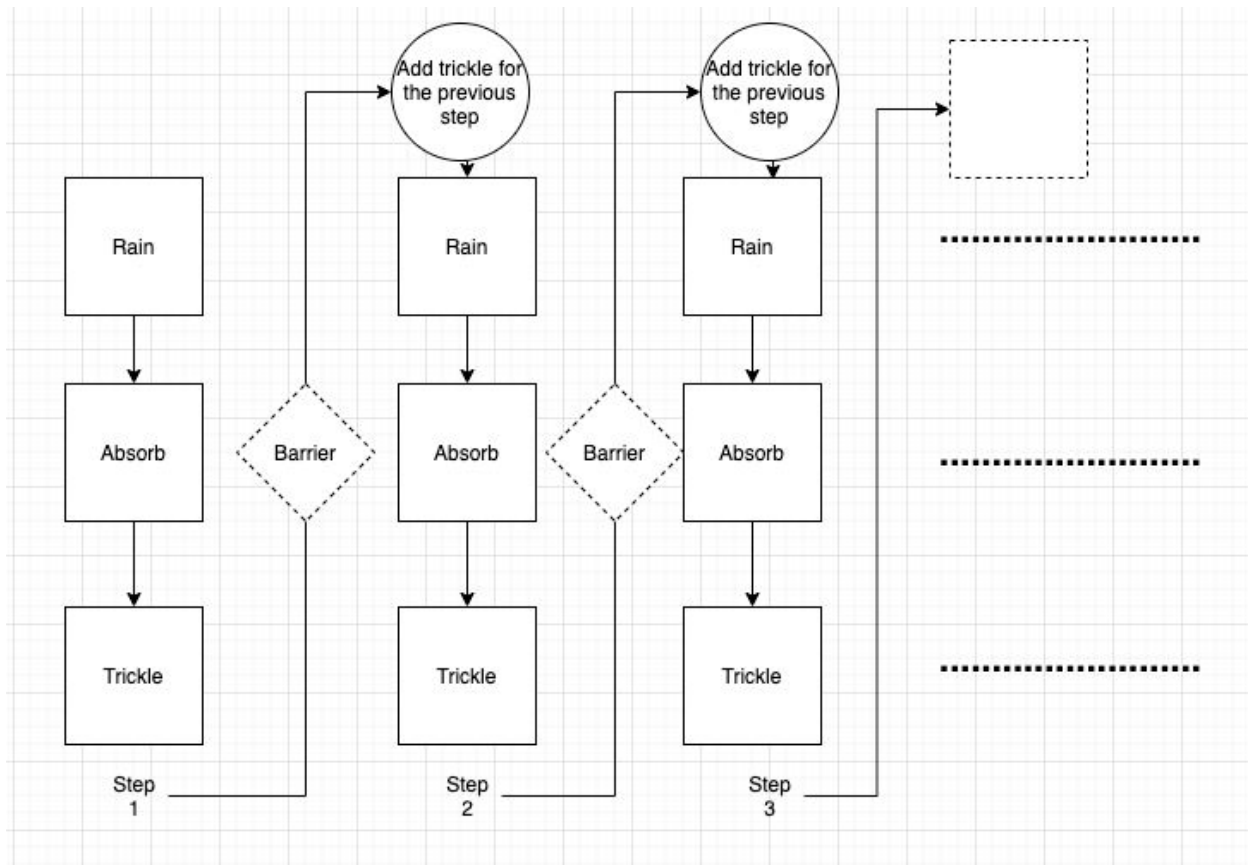
    float share = trickle[i][j] / neighborsToTrickle[i][j].size();
    rain[i][j] -= trickle[i][j];
    for (auto n : neighborsToTrickle[i][j]) {
        tempTrickle[i + direct[n[1]][0]][j + direct[n[1]][1]] += share;
    }
}

```

Parallel:

In the parallel version, we used the pre-created threads from the thread pool. And used `std::mutex` to lock the variable.

We choose to parallel the whole process of rain, absorb and trickle part in each step because that's where most of the calculation and loop happens.



`future<void> results` is used to join all the threads(works like a barrier)

```
clock_gettime(CLOCK_MONOTONIC, &start_time);

while (notDrain) {
    notDrain = false;
    future<void> results[numThreads]; // To join all the threads

    // 1) Receive a new raindrop (if it is still raining) for each point.
    // 2) If there are raindrops on a point, absorb water into the point
    // 3a) Calculate the number of raindrops that will next trickle to the
    // lowest neighbor(s)
    // cout << "-----" << wholeSteps << endl;
    for (int i = 0; i < numThreads; i++) {
        results[i] = p.push(rainAbsorbTrickle, ID++);
    }
}
```

```

    }

    for (int i = 0; i < numThreads; i++) {
        results[i].wait(); // synchronize all threads
    }

    // cout << "-----" << endl;

    ID = 0;

    nextTrickle = tempTrickle;
    tempTrickle = resetTrickle;

    --timeSteps;
    ++wholeSteps;
}

clock_gettime(CLOCK_MONOTONIC, &end_time);

```

To prevent each thread from operating on the same point when raining and absorbing, we assign each thread with one unique block of the array based on their `thread_id`:

```

void rainAbsorbTrickle(int id, int threadId) {
    for (int i = threadId * N / numThreads;
        i < threadId * N / numThreads + int(ceil(float(N) / float(numThreads)));
        ++i) {
        for (int j = (threadId * N * N / numThreads) % N;
            j < N - ((threadId * N * N / numThreads) % N + N * N / numThreads) % N;
            ++j) {
            ...
        }
    }
}

```

Since different threads operate on different blocks while raining and absorbing, there is no race condition.

The only part we need to take care of is the trickle part, because even though threads are operating on different blocks, race conditions can still happen if they write to the same neighbors when calculating the trickle, so we need to lock the **tempTrickle** array, but locking the whole array will slow down the speed greatly. After some consideration, we realize since only when the threads operate on the edge of the block can they write to the same neighbor, we only need to lock the edge of the block.

For example, if there are total $16 * 16$ points and there are 4 threads, we will assign 64 consecutive points for each thread, shown as the following

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	Thread 1															
1																
2																
3																
4	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock
5	Thread 2															
6																
7																
8																
9	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock
10	Thread 3															
11																
12																
13																
14	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock	lock
15	Thread 4															

We only need to lock the edge part which is shown in the above picture with “lock” on the point

```
void calcTrickle(int i, int j, int threadId) {
    if (trickle[i][j] == 0 || neighborsToTrickle[i][j].size() == 0) {
        return;
    }
    vector<vector<int>> direct = {
        {0, -1}, {-1, 0}, {0, 1}, {1, 0}}; // Left, up, right, down

    float share = trickle[i][j] / neighborsToTrickle[i][j].size();
    rain[i][j] -= trickle[i][j];
    for (auto n : neighborsToTrickle[i][j]) {
        if (i == threadId * N / numThreads ||
            i == threadId * N / numThreads +
                int(ceil(float(N) / float(numThreads))) - 1) {
            mtx.lock(); // Only lock on edge
        }
        tempTrickle[i + direct[n[1]][0]][j + direct[n[1]][1]] += share;
    }
}
```



```
if (i == threadId * N / numThreads ||  
    i == threadId * N / numThreads +  
        int(ceil(float(N) / float(numThreads))) - 1) {  
    mtx.unlock(); // Only lock on edge  
}  
}  
}
```

Performance:

Sequential:

```
my_4096x4096.out
1 Rainfall simulation took 1040 time steps to complete.
2 Runtime = 259.656 seconds
3
4 The following grid shows the number of raindrops absorbed at each point:
5 25 25 25 25.5 25 25 33.5 191 25 202 25 25.5 99.5
6 100 25.5 224 50 25 25 25 25.5 25.5 64 25 25 25
7 25 25 25 25 25.5 124.5 25 25 25 33.5 25 25.5 25
8 25 50 199.5 25 25 25 100 25 50 25 25.5 199 25
9 25 25.5 25 25 25 150 25 25.5 25 25 25 25 124.5
10 25 33.5 216 25.5 50 25 25 99.5 25 42 25 25 25.5
11 25 25 25 25 25 25 149.75 25 25 232.5 25 25.5 25
```

Parallel:

1 thread:

```
my_4096_1thread.out x
my_4096_1thread.out
1 Rainfall simulation took 1040 time steps to complete.
2 Runtime = 445.862 seconds
3
4 The following grid shows the number of raindrops absorbed at each point:
5 25 25 25 25.5 25 25 33.5 191 25 202 25 25.5 99.5 25
6 100 25.5 224 50 25 25 25 25.5 25.5 64 25 25 25 100
7 25 25 25 25 25.5 124.5 25 25 25 33.5 25 25.5 25 25
8 25 50 199.5 25 25 25 100 25 50 25 25.5 199 25 25
9 25 25.5 25 25 25 150 25 25.5 25 25 25 25 124.5 25
10 25 33.5 216 25.5 50 25 25 99.5 25 42 25 25 25.5 25
11 25 25 25 25 25 25 149.75 25 25 232.5 25 33.5 25 25 3
```

2 threads:

```
my_4096_2threads.out
1 Rainfall simulation took 1040 time steps to complete.
2 Runtime = 276.395 seconds
3
4 The following grid shows the number of raindrops absorbed at each point:
5 25 25 25 25.5 25 25 33.5 191 25 202 25 25.5 99.5 25 125 25 25 25 100
6 100 25.5 224 50 25 25 25 25.5 25.5 64 25 25 25 100 25 25 33.5 25 33.5
7 25 25 25 25 25.5 124.5 25 25 25 33.5 25 25.5 25 25 25 25 116.5 25 233
8 25 50 199.5 25 25 25 100 25 50 25 25.5 199 25 25 42 208 25 25 33.5
9 25 25.5 25 25 25 150 25 25.5 25 25 25 124.5 25 25 25 25 50 25 25
10 25 33.5 216 25.5 50 25 25 99.5 25 42 25 25 25.5 25 25 25 25 191.5 25
11 25 25 25 25 25 25 149.75 25 25 232.5 25 33.5 25 25 33.5 75 25 33.5 25
```

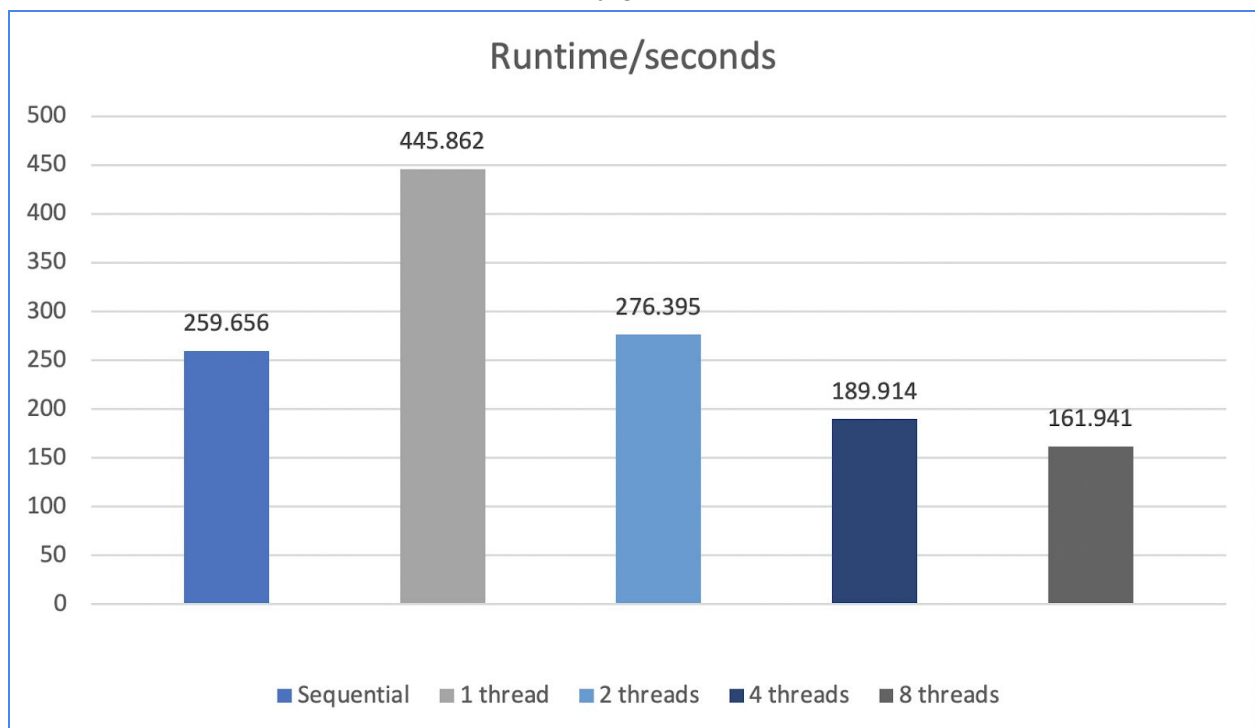
4 threads:

```
my_4096_4threads.out
1 Rainfall simulation took 1040 time steps to complete.
2 Runtime = 189.914 seconds
3
4 The following grid shows the number of raindrops absorbed at each point:
5 25 25 25 25.5 25 25 33.5 191 25 202 25 25.5 99.5 25 125 25 25 25 100
6 100 25.5 224 50 25 25 25 25.5 25.5 64 25 25 25 100 25 25 33.5 25 33.5
7 25 25 25 25 25.5 124.5 25 25 25 33.5 25 25.5 25 25 25 25 116.5 25 233
8 25 50 199.5 25 25 25 100 25 50 25 25.5 199 25 25 42 208 25 25 33.5
9 25 25.5 25 25 25 150 25 25.5 25 25 25 124.5 25 25 25 25 50 25 25
10 25 33.5 216 25.5 50 25 25 99.5 25 42 25 25 25.5 25 25 25 25 191.5 25
11 25 25 25 25 25 25 149.75 25 25 232.5 25 33.5 25 25 33.5 75 25 33.5 25
12 100 25 25 208 42 25 25.5 25 25 25.5 25 154 25 25 141.5 25 25 25 25
```

8 threads:

```
my_4096_8threads.out x
my_4096_8threads.out
1 Rainfall simulation took 1040 time steps to complete.
2 Runtime = 161.941 seconds
3
4 The following grid shows the number of raindrops absorbed at each point:
5   25   25   25   25.5   25   25   33.5   191   25   202   25   25.5   99.5   25   125   25   25   25   100
6   100   25.5   224   50   25   25   25   25.5   25.5   64   25   25   25   100   25   25   33.5   25   33.5
7   25   25   25   25   25.5   124.5   25   25   25   33.5   25   25.5   25   25   25   25   116.5   25   233
8   25   50   199.5   25   25   25   100   25   50   25   25.5   199   25   25   42   208   25   25   33.5
9   25   25.5   25   25   25   150   25   25.5   25   25   25   124.5   25   25   25   25   50   25   25
10  25   33.5   216   25.5   50   25   25   99.5   25   42   25   25   25.5   25   25   25   25   191.5   25
11  25   25   25   25   25   25   149.75   25   25   232.5   25   33.5   25   25   33.5   75   25   33.5   25
12  100   25   25   208   42   25   25.5   25   25   25.5   25   154   25   25   141.5   25   25   25   25
13  25   25   50   25   25   25   149.25   25.5   25   25   29.5   25   25   150   25   25   224.5   50   25
14  183   42   25   25   25.5   180.5   25   25   25   25   207.5   25.5   25   25   25.5   25   25   25.5   25
```

We can see that for 4 threads and 8 threads, the performance gets improved greatly. But for 2 threads and 1 thread, the performance actually gets worse.



At first, it doesn't match our expectation because performance of 1 thread is much worse and 2 threads performance is barely the same as sequential.

Then after some research, it starts to make sense to us. We think it's because of the overhead of threadpool. We can see that even when we only pre-create 1 thread in the thread pool, we still need to push the function using this single thread each step, make this single thread deal with the whole work, which will cause a lot of overhead and greatly slow down the speed. Mainly because the thread needs to reset its variables each step.

```
for (int i = 0; i < numThreads; i++) {
    results[i] = p.push(rainAbsorbTrickle, ID++);
}
```

```

for (int i = 0; i < numThreads; i++) {
    results[i].wait(); // synchronize all threads
}

```

We also used mutex in the array, the lock & unlock operation will have overhead as well. And other threads might need to wait for the current thread to finish the operation.

```

for (auto n : neighborsToTrickle[i][j]) {
    if (i == threadId * N / numThreads ||
        i == threadId * N / numThreads +
            int(ceil(float(N) / float(numThreads))) - 1) {
        mtx.lock(); // Only lock on edge
    }
    tempTrickle[i + direct[n[1]][0]][j] + direct[n[1]][1]] += share;
    if (i == threadId * N / numThreads ||
        i == threadId * N / numThreads +
            int(ceil(float(N) / float(numThreads))) - 1) {
        mtx.unlock(); // Only lock on edge
    }
}
}

```

Only when we have more threads(4,8..), the parallelization is strong enough to overcome these overhead and show improvement in the performance.